

Legionnaires' × Cooling-Tower History — Findings

2026 Upper East Side (UES) Legionnaires' disease community cluster Analysis date: 2026-07-17 · Data: NYC Open Data (Socrata) + DOH press-release list

1. Question and framing

Is there a relationship between the NYC buildings implicated in the 2026 UES Legionnaires' cluster and those buildings' cooling-tower **registration, Legionella-sampling, and inspection history**?

The DOH "clean & disinfect" list is a roster of **buildings whose cooling tower(s) tested PCR-positive for Legionella** during the July 2026 investigation — not patient addresses (those are never published). So this is a **case-control study of cooling towers**:

Outcome: building's tower(s) PCR-positive & ordered to remediate (case=1) vs not (control=0)
Exposures: historical compliance / maintenance features engineered from the two open datasets.

2. Data

Dataset	Socrata ID	Rows pulled
NYC Cooling Tower Registrations (+ Legionella sample dates)	y4fw-iqfr	5,949
NYC Cooling Tower System Inspection Results	f9wb-g8mb	122,846
DOH affected-building list (31 on 07-10 + 45 update)	press release	76 buildings

All 76 case addresses matched to a registered building (BIN) after address normalization — **0 unmatched**. (Sanity checks: 1000 Fifth Ave = 21 towers = the Met; "300 E 83rd", flagged "unregistered" by DOH, does appear in registrations.)

Leakage control: every history feature is capped at **2026-07-01** — the July-2026 investigation samples *are* the outcome, so including them would leak it.

3. Design & sample

Positivity is overwhelmingly a function of **being in the sampled area**:

Area	Buildings	Positivity
Outbreak ZIPs (10028 / 10075 / 10128)	183	40.4%
Rest of Manhattan	2,699	0.07%

Buildings outside the zone were never sampled, so counting them as "controls" would be surveillance bias. **Primary analysis is therefore restricted to the 183 registered buildings in the outbreak ZIPs (74 cases, 109 controls).**

4. Results

4a. Univariate (case vs control, in zone)

Significant / notable (Mann-Whitney or Fisher):

Feature	Case (median)	Control (median)	p
Registration age (yrs)	10.8	10.8	0.003
# Legionella samples (24 mo)	4	3	0.010
Median gap between samples (days)	41.5	54.5	0.025
Days since last sample	16	20	0.27
# Violations (all-time)	17	16	0.37
Ever cited for a violation	93%	94%	1.00
Overdue >90d on sampling	9%	18%	0.14

Direction is the opposite of a naïve "negligence" hypothesis: case towers were sampled **more** often and had **shorter** gaps, not fewer. Violation/citation history did **not** differ.

4b. Multivariable logistic regression (primary model)

183 buildings, 74 cases, 6 predictors, all VIF < 2.2. ORs are per **1 SD** (continuous). Figure: [reports/figures/odds_ratios.png](#).

Predictor	OR	95% CI	p
Registration age	1.57	0.99–2.49	0.053
# samples (24 mo)	1.49	0.94–2.35	0.089
# violations	1.13	0.78–1.63	0.51
Days since last inspection	0.81	0.58–1.12	0.20
# towers	1.00	0.65–1.54	0.99
Days since last sample	0.96	0.63–1.46	0.86

Model: McFadden $R^2 = 0.065$, in-sample AUC = 0.70, joint LLR $p = 0.013$.

L1/LASSO companion (guards against overfitting): keeps the same two leading signals — registration age (OR ≈ 1.53) and sampling frequency (OR ≈ 1.46) — shrinks # towers to zero. Cross-validated AUC ≈ 0.60 (weak).

4c. Spatial autocorrelation

Moran's I (binary weights, 400 m): case indicator **I = 0.0004**, **p = 0.15**; model residuals **I ≈ 0**, **p = 0.38**. **No fine-scale spatial clustering** of positive towers within the zone — positives and negatives are intermixed (`reports/figures/map_towers.png`).

4d. PLUTO sensitivity (building attributes)

Joining PLUTO by BBL (173/183 zone buildings matched), **case buildings are physically bigger and higher-valued**, not older:

Attribute	Case (median)	Control (median)	p
Assessed total value	\\$21.2M	\\$14.7M	0.004
Building floor area	144,280	106,034	0.025
# Floors	18	16	0.046
Building age (yrs)	62	63	0.48

Re-running the logistic model with PLUTO covariates added (`reports/odds_ratios_pluto.csv`): the **registration-age signal attenuates to null** (OR 1.20, p=0.38) — it was partly a proxy for building size/value — while **sampling frequency stays significant** (OR 1.52, p=0.023) and assessed value trends positive (OR 1.49, p=0.091). Model AUC 0.71, R² 0.072. So the most robust correlate of positivity is *how much a tower is sampled*, plausibly detection intensity, with building size/value as a secondary axis.

4e. Tier-1 additional predictors

Nine further hypotheses were tested (`src/tier1.py` , `reports/odds_ratios_tier1.csv`): severity-split violations (PHH / Critical), Legionella / water-quality-specific violations, follow-up ("non-cycle") inspections, local cooling-tower density within 200 m, sampling-cadence regularity, and ever-inactive status.

New predictor	Case vs control	p
Ever had a follow-up ("non-cycle") inspection	35% vs 20%	0.027
# non-cycle inspections	higher in cases	0.015
Water-quality / Legionella violations (focused model)	OR 1.68 / SD	0.106
Local tower density (200 m)	10 vs 10	0.60 (null)
PHH / Critical violation counts	—	0.59 / 0.24 (null)
Sampling-cadence irregularity (gap CV)	0.82 vs 0.68	0.15 (null)
Ever-inactive tower	4% vs 6%	0.74 (null)

The one nominal signal is **follow-up inspections** — complaint- or problem-triggered visits, i.e. buildings previously flagged. In the focused 6-predictor model (n=183) **sampling frequency**

remains the most robust correlate (OR 1.58, $p=0.017$); water-quality violations trend positive but are not significant; AUC 0.66.

Multiple-comparison caveat: ~9 features were tested, so ~0.5 false positives are expected at $\alpha=0.05$. The non-cycle signal ($p\approx 0.02$) does **not** survive a Bonferroni correction (~0.006) and should be read as a lead to investigate, not a finding. Tower density being null is itself informative: positive towers are **not** concentrated where towers are densest.

5. Interpretation

1. **The dominant "predictor" of a positive tower is geography, not maintenance history** — 40% positivity in-zone vs ~0% elsewhere. This is consistent with a diffuse neighborhood aerosol exposure, reinforced by the flat Moran's I (no single point-source hotspot within the zone).
2. **Within the outbreak zone, compliance history explains little** (AUC 0.70, R^2 0.07). Crucially, the metrics a regulator would flag — violations, citations, overdue sampling, sparse inspections — were **not** associated with testing positive. This argues the cluster is **not** simply a "negligent operators" story.
3. The two weak positive associations both plausibly reflect **detection/confounding, not causation**: towers sampled more frequently are more likely to *catch* a positive (surveillance intensity), and older registrations proxy for older/larger building systems. Neither is evidence that more testing *causes* Legionella.

Bottom line: there is a modest, statistically borderline relationship between a building's cooling-tower history and outbreak positivity, but it is weak, points opposite to a negligence hypothesis, and is dwarfed by location. The data do not support singling out poorly-maintained towers as the driver of this cluster.

6. Limitations

- Outcome = **tower positivity**, not human infection (patient addresses private).
- Small case count (74 in-zone) → wide CIs; borderline p-values are fragile.
- **Surveillance/detection bias:** positives were actively *sought* in these ZIPs.
- ~6-month posting lag on inspections; owner-reported sample dates may be imperfect.
- Registration age is a proxy (date first registered $\approx$ 2015 floor), not true tower age.
- Association $\neq$ causation; ecological/aggregation caveats apply.

7. Reproducibility

```
python3 src/fetch_data.py           # pull Socrata datasets -> data/raw/
python3 src/build_case_list.py      # 76-building case roster -> data/raw/
python3 src/match_and_label.py      # address match/audit -> data/processed/
python3 src/build_features.py       # labeled feature table -> data/processed/
python3 src/analyze.py              # EDA + regression + spatial -> reports/
python3 src/enrich_pluto.py         # PLUTO join + sensitivity model -> reports/
python3 reports/build_map.py        # self-contained artifact map (NYC grid basemap)
python3 reports/build_leaflet.py    # standalone map w/ OpenStreetMap tiles (open locally)
```

Console logs: `reports/stats.txt` , `reports/stats_pluto.txt` . Odds ratios: `reports/odds_ratios.csv` , `reports/odds_ratios_pluto.csv` .

8. Interactive maps

- `reports/map.html` — self-contained interactive dot-map over the authentic NYC street grid (basemap from NYC Open Data `inkn-q76z`); hover history, filters, table view; works inside a sandboxed artifact (no external requests).
- `reports/map_leaflet.html` — standalone map on a live OpenStreetMap tile basemap; open in a browser locally (loads external tiles, so not for the sandbox).

When the DOH list grows further, drop the new addresses into `src/build_case_list.py` and re-run — the pipeline is a drop-in re-run.